

ECE 228 Spring 2025 Lecture 13: Neural Operators Part 1

Luke Bhan

May 13, 2025

Neural ODEs introduced us to this idea of learning continuous representations. How can we take this further?

Recall that Neural ODEs learn the flow field f of a dynamical system

$$\frac{dz(t)}{dt} = f(z(t), t; \theta), \quad z(0) = \bar{z}_0.$$

This is the first type of continuous learning representation we have discussed. However, there's more to be explored here. NeuralODEs require us to build a set of points $t \in [0, T]$ to learn the solution of which we use the final output time T .

Motivating question: How can we design neural architectures that we can query at any time (or for PDEs, any space) location?

First a little history: How do we simulate PDEs (See homework 3)

① **Start with a PDE:**

$$\frac{\partial u}{\partial t} = \mathcal{F}(u, \nabla u, \nabla^2 u, \dots)$$

This represents the evolution of a physical quantity u .

② **Discretize Space and Time:** Replace continuous space/time with a grid:

$$x_0, x_1, \dots, x_N \quad \text{and} \quad t_0, t_1, \dots, t_M$$

③ **Approximate Derivatives:** Use finite differences, finite volumes, or finite elements to approximate spatial and temporal derivatives.

④ **Step Forward in Time:** Update the solution using a time-stepping method (e.g., Euler, Runge-Kutta).

⑤ **Repeat:** Continue iterating until the final time is reached.

History: Key limitation of numerically simulating PDEs




Figure: Example of a triangle mesh for a PDE.

- **Fixed Resolution:** Simulations are restricted to predefined spatial and temporal grid points. Fine grids increase accuracy but also computational cost and can lead to millions of simulation points
- **Poor Generalization:** Traditional solvers compute solutions only at grid points they're run on. New parameters require *re-running the simulation from scratch*.

How do we *classically* solve PDEs?

Take the simplest PDE (1D transport):

$$\begin{aligned}\frac{\partial u(x, t)}{\partial t} &= \frac{\partial u(x, t)}{\partial x}, \quad x \in (0, 1), t \in \mathbb{R}^+, \\ u(x, 0) &= f(x), \\ u(t, 0) &= 0.\end{aligned}$$

This PDE just shifts the initial condition with space and time.

Example of a transport PDE




Figure: Transport PDE with Gaussian initial condition centered at $x = 0.9$.

How do we *classically* solve PDEs?

Can apply Euler's method:

$$\begin{aligned}\frac{\partial u(x, t)}{\partial t} &\approx \frac{u(x, t + \delta t) - u(x, t)}{\delta t} + \mathcal{O}(\delta t^2), \\ \frac{\partial u(x, t)}{\partial x} &\approx \frac{u(x + \delta x, t) - u(x, t)}{\delta x} + \mathcal{O}(\delta x^2),\end{aligned}$$

Combining yields the **finite difference** scheme

$$u(x, t + \delta t) = u(x, t) + \frac{\delta t}{\delta x} (u(x + \delta x, t) - u(x, t))$$

where we impose the conditions

$$\begin{aligned}u(x, 0) &= f(x), \\ u(t, 0) &= 0.\end{aligned}$$

Then, to simulate until say $T = 5$, we just recursively run the above calculation until our desired timestep is reached.

PDE solvers get very expensive very quickly

Unlike Euler's method for ODEs, PDE solvers become computationally costly extremely quickly.

- This is due to the fact we have both a spatial and temporal component, and thus we need to solve this PDE across both space and time.
- Furthermore, As with ODE's we require the step sizes δx and δt to be small. A reasonable choice is $\delta x = 0.01$ and $\delta t = 0.0005$ (for most numerical scheme $\delta t \ll \delta x$ as discussed in your homework).
- Then, to solve to $T = 5$, we require $(1/\delta x + 1) \times (5/\delta t + 1) = 1,010,101$ points of evaluation.

Solving computationally becomes much worse as dimension scales

Consider the Navier Stokes PDE in 2D.

- For 1D, to solve to $T = 5$, we require $(1/\delta x + 1) \times (5/\delta t + 1) = 1,010,101$ points of evaluation.
- For 2D, to solve to $T = 5$, we require $(1/\delta y + 1) \times (1/\delta x + 1) \times (5/\delta t + 1) = 10,211,201$ points.
- For 3D, for $T = 5$, we require a billion points $(1,031,331,301)$ which is not possible on everyday machines.



Figure: Sierra supercomputer in Livermore California.

How can we design a machine learning model that can learn PDE solutions and be queried at any grid location?

Answer: Design a machine learning model using *continuous representations*.

To do so, we need to introduce the mathematical notion of an **operator**. One reasonable viewpoint is too look at an operator as an abstraction of functions. That is, functions map vectors to vectors:




Figure: The function f take an input vector $x \in \mathbb{R}^n$ and maps to an output vector in \mathbb{R}^m .

What is an operator?

Now, imagine we considered a set of functions, for example all polynomials of degree 2 or all sin function of the form $f(x) = A\sin(Bx + C)$ where A, B, C can vary. Then, like a function that maps vectors to vectors, an operator maps *functions to functions*:





Figure: Let \mathcal{U} be a set of functions and \mathcal{V} be another set of functions. Then the operator \mathcal{G} takes a function in $u \in \mathcal{U}$ and maps it to a function $v \in \mathcal{V}$.

Examples of operators

You all (Y'all for the southerners) know many types of operators. Perhaps, the most famous are:

- The matrix multiplication A is *linear* operator:



- The integral operator. Let $f \in \mathcal{F}$ and $g \in \mathcal{G}$, then the integral acts as

$$g = \int f(x) dx.$$

- Analogously the derivative of a function is an operator.
- Another operator is the shift operator. Given a parameter c , it acts as:


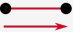
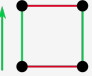
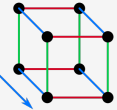
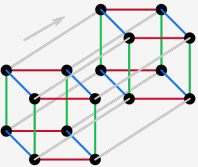

$$g(x) = \mathcal{G}[f](x) = f(x + c)$$

An important property of vector spaces

In vector spaces, we had this important structure of dimension which essentially told us:

- We need exactly d linearly independent vectors to create any vector in the space (recall this is a basis)

This essentially means that our vector field has d degrees of information or one can think of this as d values need to specify the object.

					<div><div>X</div><div>Y</div><div>Z</div><div>W</div></div>
0	1	2	3	4	#Dim

Operators typically make function spaces, which in many cases have ∞ dimension

Consider the set of all continuous functions \mathcal{C} .

Using the fact that dimension says we need d number of linearly independent objects to build any object in the space, what dimension do you think \mathcal{C} is?

Answer: \mathcal{C} is infinite dimensional! This will be an important property when we consider operator mappings, as we *cannot* write down an infinite dimensional object in the form of a finite number of basis vectors! This will become important in our neural network design.

Approximation of operators

Neural networks approximate functions. Neural operators approximate operators.

What does it mean to approximate an operator?

Main questions:

- How do we input a function f which is inherently continuous and (sometimes) a member of an infinite dimensional set into a computer?
- Say we can do this, what does it mean to get an output function from our neural network?
- Why would anyone want to do this?

Let's answer the third question before moving on. Do you see any advantages?

Why do we care about approximating operators?

The solution to dynamical systems - both ODEs and PDEs is inherently an *operator*. For example consider ODEs

$$\frac{dz(t)}{dt} = f(z(t), t), \quad z(0) = \bar{z}_0.$$

Then, one can write the solution as

$$z(t) = \int_0^t f(z(t), t) dt + \bar{z}_0.$$

The solution to an ODE is given by an **operator**!

Thus, imagine we want to learn $z(t)$ for any input \bar{z}_0 . Then, if we can learn this implicit integral operator, we have a solution to the ODE for all times t with all initial conditions \bar{z}_0 removing the expensive numerical schemes (Euler's, RK4, etc.)

Perhaps you think ODEs are too simple

Consider the holy grail of PDE's: Navier-Stokes

$$\rho \left(\frac{\partial \mathbf{u}}{\partial t} + \mathbf{u} \cdot \nabla \mathbf{u} \right) = -\nabla p + \mu \nabla^2 \mathbf{u} + \mathbf{f},$$
$$\nabla \cdot \mathbf{u} = 0,$$

where:

- \mathbf{u} is the velocity field
- p is the pressure field
- ρ is the fluid density (constant for incompressible flow)
- μ is the dynamic viscosity
- \mathbf{f} represents body forces (e.g., gravity)

First, if you can just prove that a solution exists congrats your a millionaire.^a

^ahttps://en.wikipedia.org/wiki/Millennium_Prize_Problems

Operator formulation of Navier-Stokes

Given an initial velocity field at time 0, $\mathbf{u}(\mathbf{x}, 0) = \mathbf{g}(\mathbf{x})$, $\mathbf{x} \in \mathbb{R}^n$, then we can write the solution operator as

$$\mathbf{u}(\mathbf{x}, t) = \mathcal{G}[p, \rho, \mu, \mathbf{f}, \mathbf{g}](\mathbf{x}, t)$$

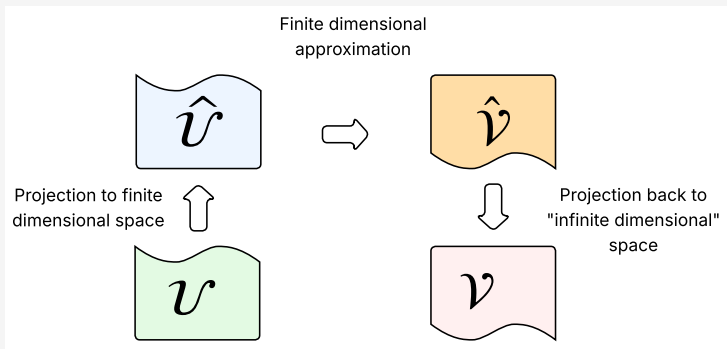
which takes an input of $\mathbb{R} \times \mathbb{R} \times \mathbb{R} \times C^0(\mathbb{R}^n \times \mathbb{R}^0; \mathbb{R}^n) \times C^0(\mathbb{R}^n; \mathbb{R}^n)$ to an output of some function \mathbf{u} (of which we don't know if it's smooth. See Millennium problems).

- Notice, for PDEs, the operator learned is a mapping of function spaces.
- As with the ODE problem, if we can learn an approximation of this operator, then we obtain the Navier-Stokes equations at any desired \mathbf{x}, t . The grid independence or otherwise known as **discretization invariance** will enable us to perform calculations that are otherwise impossible with traditional solvers.

So why learn an approximation of an operator

It gives us a framework to develop machine learning solutions to ODE and PDEs that can be queried at any desired x, t location!

So what does this look like in practice? **Answer:** There are a variety of approaches, but they all follow the same structure.



Let's walk through an example: Transport PDE and FNO

Consider the 1D transport PDE:

$$\begin{aligned}\frac{\partial u(x, t)}{\partial t} &= \frac{\partial u(x, t)}{\partial x}, \quad x \in (0, 1], t \in \mathbb{R}^+, \\ u(x, 0) &= f(x), \\ u(t, 0) &= 0.\end{aligned}$$

Given, $f(x)$, we want to know the solution of the PDE from time $T = 3$ to $T = 5$. That is, we want to learn the operator mapping

$$f(x) \mapsto u(x, t), \quad t \in [3, 5], x \in [0, 1].$$

Step 1: How do we encode $f(x)$ (project $f(x)$ to finite dimensional space)?

Encoding: Simplest approach.

That is, consider points $x_1, x_2, \dots, x_n \in [0, 1] = X$. One possible encoding of $f(x)$ is to evaluate the function at these points:

$$[f(x_1), f(x_2), \dots, f(x_n)].$$

More commonly, one extends this by a classical feedforward neural network with $[f(x_1), \dots, f(x_n)]$ as its input (sometimes, this projects to a *higher dimensional space* - you can think of the NN as performing interpolation between the x_k points).

Key: This introduces error in our approximation as we can't pass the true functional form, but this is a representation which can be passed into a computer as a vector.

How to do a finite dimensional approximation of an operator

Let a typical MLP for a neural network be given by

$$f(x) = \sigma(Wx + b).$$

Then, a neural operator approximation is given by

$$\mathcal{L}(f)(x) := \sigma(Wf(x) + b + (\mathcal{K}_I f)(x))$$

where

$$(\mathcal{K}_I f)(x) = \int_{y \in X} K_I(x, y) f(y) dy.$$

So what is new: Essentially, we added this **non-local** integral term which now depends on every point in the space and multiplies by the weight matrix K_I .

Why this structure for the hidden layers?

Two reasons:


- Theoretically, this structure has the nice *universal operator approximation* properties that we had for neural networks and functions
- In practice, it yields a very rich set of architectures based on the choice of K_l which achieve state-of-the-art results.

Can you think of any connection or differences to other architectures we have discussed (Transformers, ResNets, CNNs)?

Decoder structure

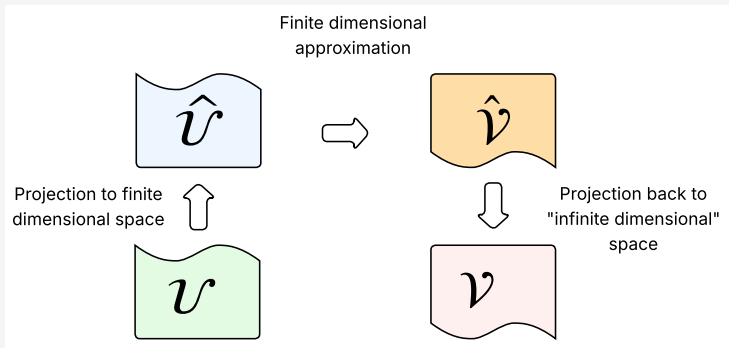
Take the output of the hidden layers and *project* back to the desired evaluation points.

How this projection is done can vary, and there are a variety of approaches (Example of one approach later today).




Theory recap: Neural Operators

- ① **Encoder:** Obtain some point-wise representation of our functional input.
- ② **Hidden layers:** Perform *Non-local* finite dimensional approximation with our kernel matrix
- ③ **Decoder:** Project back to desired grid locations (architecture dependent)



Architecture Example: Deep Operator Networks (DeepONet)

- 1 Encoder: Standard Point-wise function evaluations
 $[f(x_1), \dots, f(x_n)]$
- 2 Finite dimensional approximation: Let us consider p different neural networks (MLP) that all take in $[f(x_1), \dots, f(x_n)]$ as their input and output β_1, \dots, β_p .
- 3 Decoder: To construct our target function v , we use a second NN that takes in the evaluation point y which is where to evaluate the target function and perform a linear combination.



Architecture Intuition: Deep Operator Networks (DeepONet)

Main idea: Use the finite dimensional approximation to learn a *basis* for the target function which can be queried at any point.

That is, from the functional inputs $f(x_1), \dots, f(x_n)$, learn basis vectors b_1, \dots, b_p .

Then, to decode, pass in the target location y and learn the coefficients t_1, \dots, t_p that act as the linear combination coefficients for b_1, \dots, b_p to obtain the result:

$$\mathcal{G}(f)(y) \approx \sum_{i=1}^p b_i([f(x_1), \dots, f(x_n)]) \times \tau_i(y)$$

- b_i are given via neural networks (MLP, **Branch net**).
- τ_i are given via neural networks (MLP, **Trunk net**)

Connection to the theoretical definition of Neural Operators

Think of the summation as the kernel component of the DeepONet:

$$\mathcal{G}(f)(y) \approx \sum_{i=1}^p b_i([f(x_1), \dots, f(x_n)]) \times \tau_i(y)$$

$$\mathcal{L}(f)(x) := \sigma(Wf(x) + b + (\mathcal{K}f)(x))$$

That is, treat $\sum_{i=1}^p b_i(f) \times \tau_i(y)$ as the kernel function and let $W, b = 0$ and simplify $\sigma = \text{Identity}$. Then the DeepONet is a simplified form of the abstract Neural Operator definition. ¹

¹See Section 3.6 of <https://arxiv.org/pdf/2304.13221> for more detail.

Key features of DeepONet

- Can query the NN at any desired location y *without retraining* (Super-resolution). Branch and trunk NNs are independently computable.
- Easily extended where the two networks are of any form.
- Theoretical universal approximation guarantees.
- "Notion of interpretability" for the basis vectors.

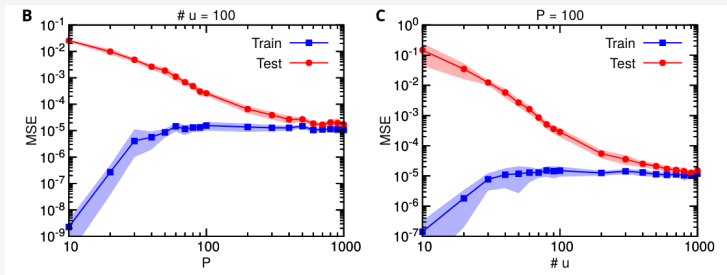


Figure: (Left) scaling with respect to number of evaluation points x_1, \dots, x_P . (Right) Scaling with respect to number of training data. Both for a reaction diffusion PDE.

What can you think of?

- The linear combination structure is less expressive Do these types of dot products have issues in other architectures?
- Can struggle with generalization and be very sample inefficient due to its very generic architecture.
- Resolution mismatch - if your training data is given at uniform grid, will struggle to evaluate on non-uniform grids.

Extensions of DeepONet

There have been a variety of extensions enhancing DeepONet.

- NOMAD - Generalizes the dot-product in DeepONet to be the input to a third NN (Paper link).
- Use DeepONet with pretrained encoder-decoder structure of functional evaluations for better encoding/decoding. (Paper link)
- Multiscale DeepONet that solves the non-uniform grid challenge (Paper link).
- Integration with RNNs for temporal dependence (Paper link)
- Uncertainty quantification ingrained DeepONets (Paper link)

Next time...

- Introduce other types of neural operator architectures (Fourier Neural Operator)




Figure: Fourier Neural Operator Architecture

- Discuss whether super-resolution really works?
- If time permits, a few examples of how these can be applied in climate modeling, robotics, and earthquakes (See your homework).